

SQL Review

Create Table with Foreign Keys

```
CREATE TABLE Studios(  
    name VARCHAR(20) PRIMARY KEY,  
    website VARCHAR(255)  
);
```

```
CREATE TABLE Stars (  
    name VARCHAR(20) PRIMARY KEY,  
    gender CHAR(1),  
    birthyear INT,  
    birthplace VARCHAR(40)  
);
```

```
CREATE TABLE Movies (  
    title VARCHAR(50),  
    year INT,  
    length INT,  
    rating CHAR(2),  
    studioname VARCHAR(20) REFERENCES Studios(name) ON DELETE CASCADE,  
    PRIMARY KEY (title, year)  
);
```

```
CREATE TABLE StarsIn (  
    title VARCHAR(50),  
    year INT,  
    starname VARCHAR(20),  
    PRIMARY KEY (title, year, starname),  
    FOREIGN KEY (title, year) REFERENCES Movies(title, year) ON DELETE CASCADE,  
    FOREIGN KEY (starName) REFERENCES Stars(name) ON DELETE CASCADE  
);
```

Insert – Studios

```
INSERT INTO Studios  
VALUES ('Fox', 'foxmovies.com');
```

```
INSERT INTO Studios  
VALUES ('Disney', 'disney.com');
```

```
INSERT INTO Studios  
VALUES ('Paramount', 'www.paramount.com');
```

Insert – Movies

```
INSERT INTO Movies
```

```
VALUES('Walk the Line', 2005, 136, 'PG', 'Fox');
```

```
INSERT INTO Movies
```

```
VALUES('Pretty Woman', 1990, 119, 'R', 'Disney');
```

```
INSERT INTO Movies
```

```
VALUES('Wayne's World', 1991, 104, 'PG', 'Paramount');
```

```
INSERT INTO Movies
```

```
VALUES('Unfaithful', 2002, 124, 'R', 'Fox');
```

```
INSERT INTO Movies
```

```
VALUES('Runaway Bride', 1999, 116, 'PG', 'Paramount');
```

```
INSERT INTO Movies
```

```
VALUES('The Princess and the Frog', 2009, 97, 'G', 'Disney');
```

Insert – Stars

```
INSERT INTO Stars  
VALUES('Richard Gere', 'M', 1949, 'Philadelphia, Pennsylvania, USA');
```

```
INSERT INTO Stars  
VALUES('Joaquin Phoenix', 'M', 1974, 'San Juan, Puerto Rico');
```

```
INSERT INTO Stars  
VALUES('Reese Witherspoon', 'F', 1976, 'Baton Rouge, Louisiana, USA');
```

```
INSERT INTO Stars  
VALUES('Julia Roberts', 'F', 1967, 'Smyrna, Georgia, USA');
```

```
INSERT INTO Stars  
VALUES('Mike Myers', 'M', 1963, 'Scarborough, Ontario, Canada');
```

```
INSERT INTO Stars  
VALUES('Oprah Winfrey', 'F', 1954, 'Kosciusko, Mississippi, USA');
```

Insert – StarsIn

```
INSERT INTO StarsIn VALUES('Walk the Line', 2005, 'Joaquin Phoenix');
```

```
INSERT INTO StarsIn VALUES('Walk the Line', 2005, 'Reese Witherspoon');
```

```
INSERT INTO StarsIn VALUES('Pretty Woman', 1990, 'Richard Gere');
```

```
INSERT INTO StarsIn VALUES('Pretty Woman', 1990, 'Julia Roberts');
```

```
INSERT INTO StarsIn VALUES('Wayne's World', 1991, 'Mike Myers');
```

```
INSERT INTO StarsIn VALUES('Unfaithful', 2002, 'Richard Gere');
```

```
INSERT INTO StarsIn VALUES('Runaway Bride', 1999, 'Richard Gere');
```

```
INSERT INTO StarsIn VALUES('Runaway Bride', 1999, 'Julia Roberts');
```

```
INSERT INTO StarsIn VALUES('The Princess and the Frog', 2009, 'Oprah Winfrey');
```

Creation and insertion order

1. **Movies** after **Studios**
2. **StarsIn** after **Movies** and **Stars**

Ordering the Input

Example. Find the Disney movies and list them by length, shortest first.

```
SELECT *  
FROM Movies  
WHERE studioName = 'Disney'  
ORDER BY length;
```

Example. Find the Disney movies and list them by length, shortest first, and among movies of equal length, sort alphabetically.

```
SELECT *  
FROM Movies  
WHERE studioName = 'Disney'  
ORDER BY length, title;
```

Remarks

- Ordering is ascending, unless you specify the DESC keyword after **an attribute**.
- Ties are broken by the second attribute on the ORDER BY list, etc.

Natural Join with USING

Better than NATURAL JOIN:

```
SELECT *  
FROM Movies JOIN StarsIn USING (title,year);
```

Because now it is explicit which attributes are used to join the tables.

Join with ON

A similar result can be obtained by:

```
SELECT *  
FROM Movies JOIN StarsIn ON  
        Movies.title=StarsIn.title AND  
        Movies.year=StarsIn.year;
```

Exactly same as:

```
SELECT *  
FROM Movies, StarsIn  
WHERE    Movies.title=StarsIn.title AND  
        Movies.year=StarsIn.year;
```

Outer Joins

```
SELECT *  
FROM Movies FULL OUTER JOIN StarsIn USING(title,year);
```

One of LEFT, RIGHT, or FULL before OUTER (but not missing).

- ◆ LEFT = pad dangling tuples of Movies only.
- ◆ RIGHT = pad dangling tuples of StarsIn only.
- ◆ FULL = pad both.

Union/Intersection/Difference

```
SELECT title, year  
FROM StarsIn  
WHERE starName='Richard Gere'
```

UNION / INTERSECT / EXCEPT (use one of them depending on request)

```
SELECT title, year  
FROM StarsIn  
WHERE starName='Julia Roberts' ;
```

Aliases

Find pairs of stars who have played together in the same movie.

```
SELECT S1.starname, S2.starname
FROM StarsIn S1, StarsIn S2
WHERE S1.title = S2.title AND S1.year = S2.year
      AND S1.starname < S2.starname;
```

Grouping

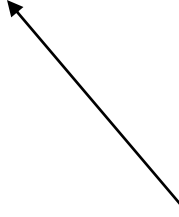
```
SELECT studioName, AVG(length)
FROM Movies
GROUP BY studioName;
```

Another Example

From **Movies** and **StarsIn**, find the star's total length of film played.

```
SELECT starName, SUM(length)
FROM Movies, StarsIn
WHERE Movies.title=StarsIn.title AND Movies.year=StarsIn.year
GROUP BY starName;
```

Compute
those
tuples first,
then group
by starName.



HAVING Clauses

```
SELECT starName, SUM(length)
FROM Movies, StarsIn
WHERE Movies.title=StarsIn.title AND Movies.year=StarsIn.year
GROUP BY starName
HAVING MIN(StarsIn.year) < 2000;
```


Requirements on HAVING Conditions

- They may refer to attributes that make sense within a group; i.e., they are either:
 1. Grouping attribute, or
 2. Aggregated attribute.

Restriction on SELECT Lists With Aggregation

- If any aggregation is used, then each element of the SELECT list must be either:
 1. Aggregated, or
 2. An attribute on the GROUP BY list.

Exercise

Using **Movies**, **StarsIn**, and **Stars**,

find the star's total length of film played.

We are interested only in Canadian stars and
who first appeared in a movie before 2000.

```
SELECT starName, SUM(length)
FROM Movies, StarsIn, Stars
WHERE Movies.title=StarsIn.title AND Movies.year=StarsIn.year
      AND Stars.name=StarsIn.starName
      AND Stars.birthplace LIKE '%Canada%'
GROUP BY starName
HAVING MIN(StarsIn.year) < 2000;
```

Correlated Subqueries

- Suppose StarsIn table has an additional attribute “salary”
StarsIn(movie, movie, starName, salary)

Now, find the stars who were paid for some movie more than the average salary for that movie.

```
SELECT starName, title, year
FROM StarsIn X
WHERE salary >
      (SELECT AVG(salary)
       FROM StarsIn
       WHERE title = X.title AND year=X.year);
```

Remark

Semantically, the value of the X tuple changes in the outer query, so the database must rerun the subquery for each X tuple.

Another Solution (Nesting in FROM)

```
SELECT X.starName, X.title, X.year
FROM StarsIn X, (SELECT title, year, AVG(salary) AS avgSalary
                  FROM StarsIn
                  GROUP BY title, year) Y
WHERE X.salary>Y.avgSalary AND
      X.title=Y.title AND X.year=Y.year;
```

Views

- A view is a “*virtual table*”, a relation that is defined in terms of the contents of other tables and views.

Example

```
CREATE VIEW DMovies AS  
  SELECT title, year, length, rating  
  FROM Movies  
  WHERE studioName = 'Disney';
```

Constraints – mutually exclusive subclasses

```
CREATE TABLE Vehicles (  
    vin CHAR(17) PRIMARY KEY,  
    vehicle_type CHAR(3) CHECK(vehicle_type IN ('SUV', 'ATV')),  
    fuel_type CHAR(4),  
    door_count INT CHECK(door_count >= 0),  
    UNIQUE(vin, vehicle_type)  
);
```

```
CREATE TABLE SUVs (  
    vin CHAR(17) PRIMARY KEY,  
    vehicle_type CHAR(3) CHECK(vehicle_type = 'SUV'),  
    FOREIGN KEY (vin, vehicle_type) REFERENCES Vehicles (vin, vehicle_type)  
);
```

```
CREATE TABLE ATVs (  
    vin CHAR(17) PRIMARY KEY,  
    vehicle_type CHAR(3) CHECK(vehicle_type = 'ATV'),  
    FOREIGN KEY (vin, vehicle_type) REFERENCES Vehicles (vin, vehicle_type)  
);
```

```
CREATE VIEW ParamountMovies AS
  SELECT title, year
  FROM Movies
  WHERE studioName = 'Paramount'
WITH CHECK OPTION;
```

```
INSERT INTO ParamountMovies VALUES ('Star Trek', 1979);
```

Insertion fails!

Why?

Rationale for this behavior is:

- Were the insertion allowed, it would insert a tuple with NULL for studioName in base table Movie.
- However, such a tuple doesn't satisfy the condition for being in the ParamountMovie view! (“**invisible tuple to the view**”)
- Thus, it shouldn't be allowed to get into the database through the ParamountMovie view.

Another example: Cars

-- Another example: Cars

```
CREATE TABLE cars (  
  car_id serial primary key,  
  car_name varchar(255),  
  brand VARCHAR(20)  
);
```

```
CREATE VIEW audi_cars AS  
  SELECT car_id, car_name, brand  
  FROM cars  
  WHERE brand = 'Audi'  
WITH CHECK OPTION;
```

```
CREATE VIEW ford_cars AS  
  SELECT car_id, car_name, brand  
  FROM cars  
  WHERE brand = 'Ford'  
WITH CHECK OPTION;
```

```
INSERT INTO audi_cars (car_name,brand) VALUES('Q2','Audi');  
INSERT INTO audi_cars (car_name,brand) VALUES('S1','Audi');
```

```
INSERT INTO ford_cars (car_name,brand) VALUES('Edge','Ford');  
INSERT INTO ford_cars (car_name,brand) VALUES('Mustang','Ford');
```

```
INSERT INTO ford_cars(car_name,brand) VALUES('RS6 Avant','Audi');  
-- ERROR: new row violates check option for view "ford_cars"
```